

Join Data with



Details



⚠ The class of service you searched may not be available on one or more flights

BNA - ORD

Flight 1 of 2

ORD - YVR

Flight 2 of 2

Nashville, TN to Chicago, IL

Thursday, July 26, 2018

4:10 PM → 6:03 PM

AA 3246 ■ CRJ-900 RJ 700
Operated by SkyWest Airlines As American Eagle

Travel info

Travel time: 1h 53m
Connection time: 2h 33m

Performance

On time: 52%
Late: 43%

Performance*

On time: 52%**
Late: 43%

Main Cabin

Meals: Beverage service
Booking code: V
Class: Economy

Business

Meals: Beverage service
Booking code: I
Class: First

* This is based on information from the month of May 2018

** The on-time arrival percentage for the selected flight is based on arrival within 14 minutes after

**** The on-time arrival percentage for the selected flight is based on arrival within 14 minutes after the scheduled arrival as reported monthly to the U.S. Department of Transportation.**

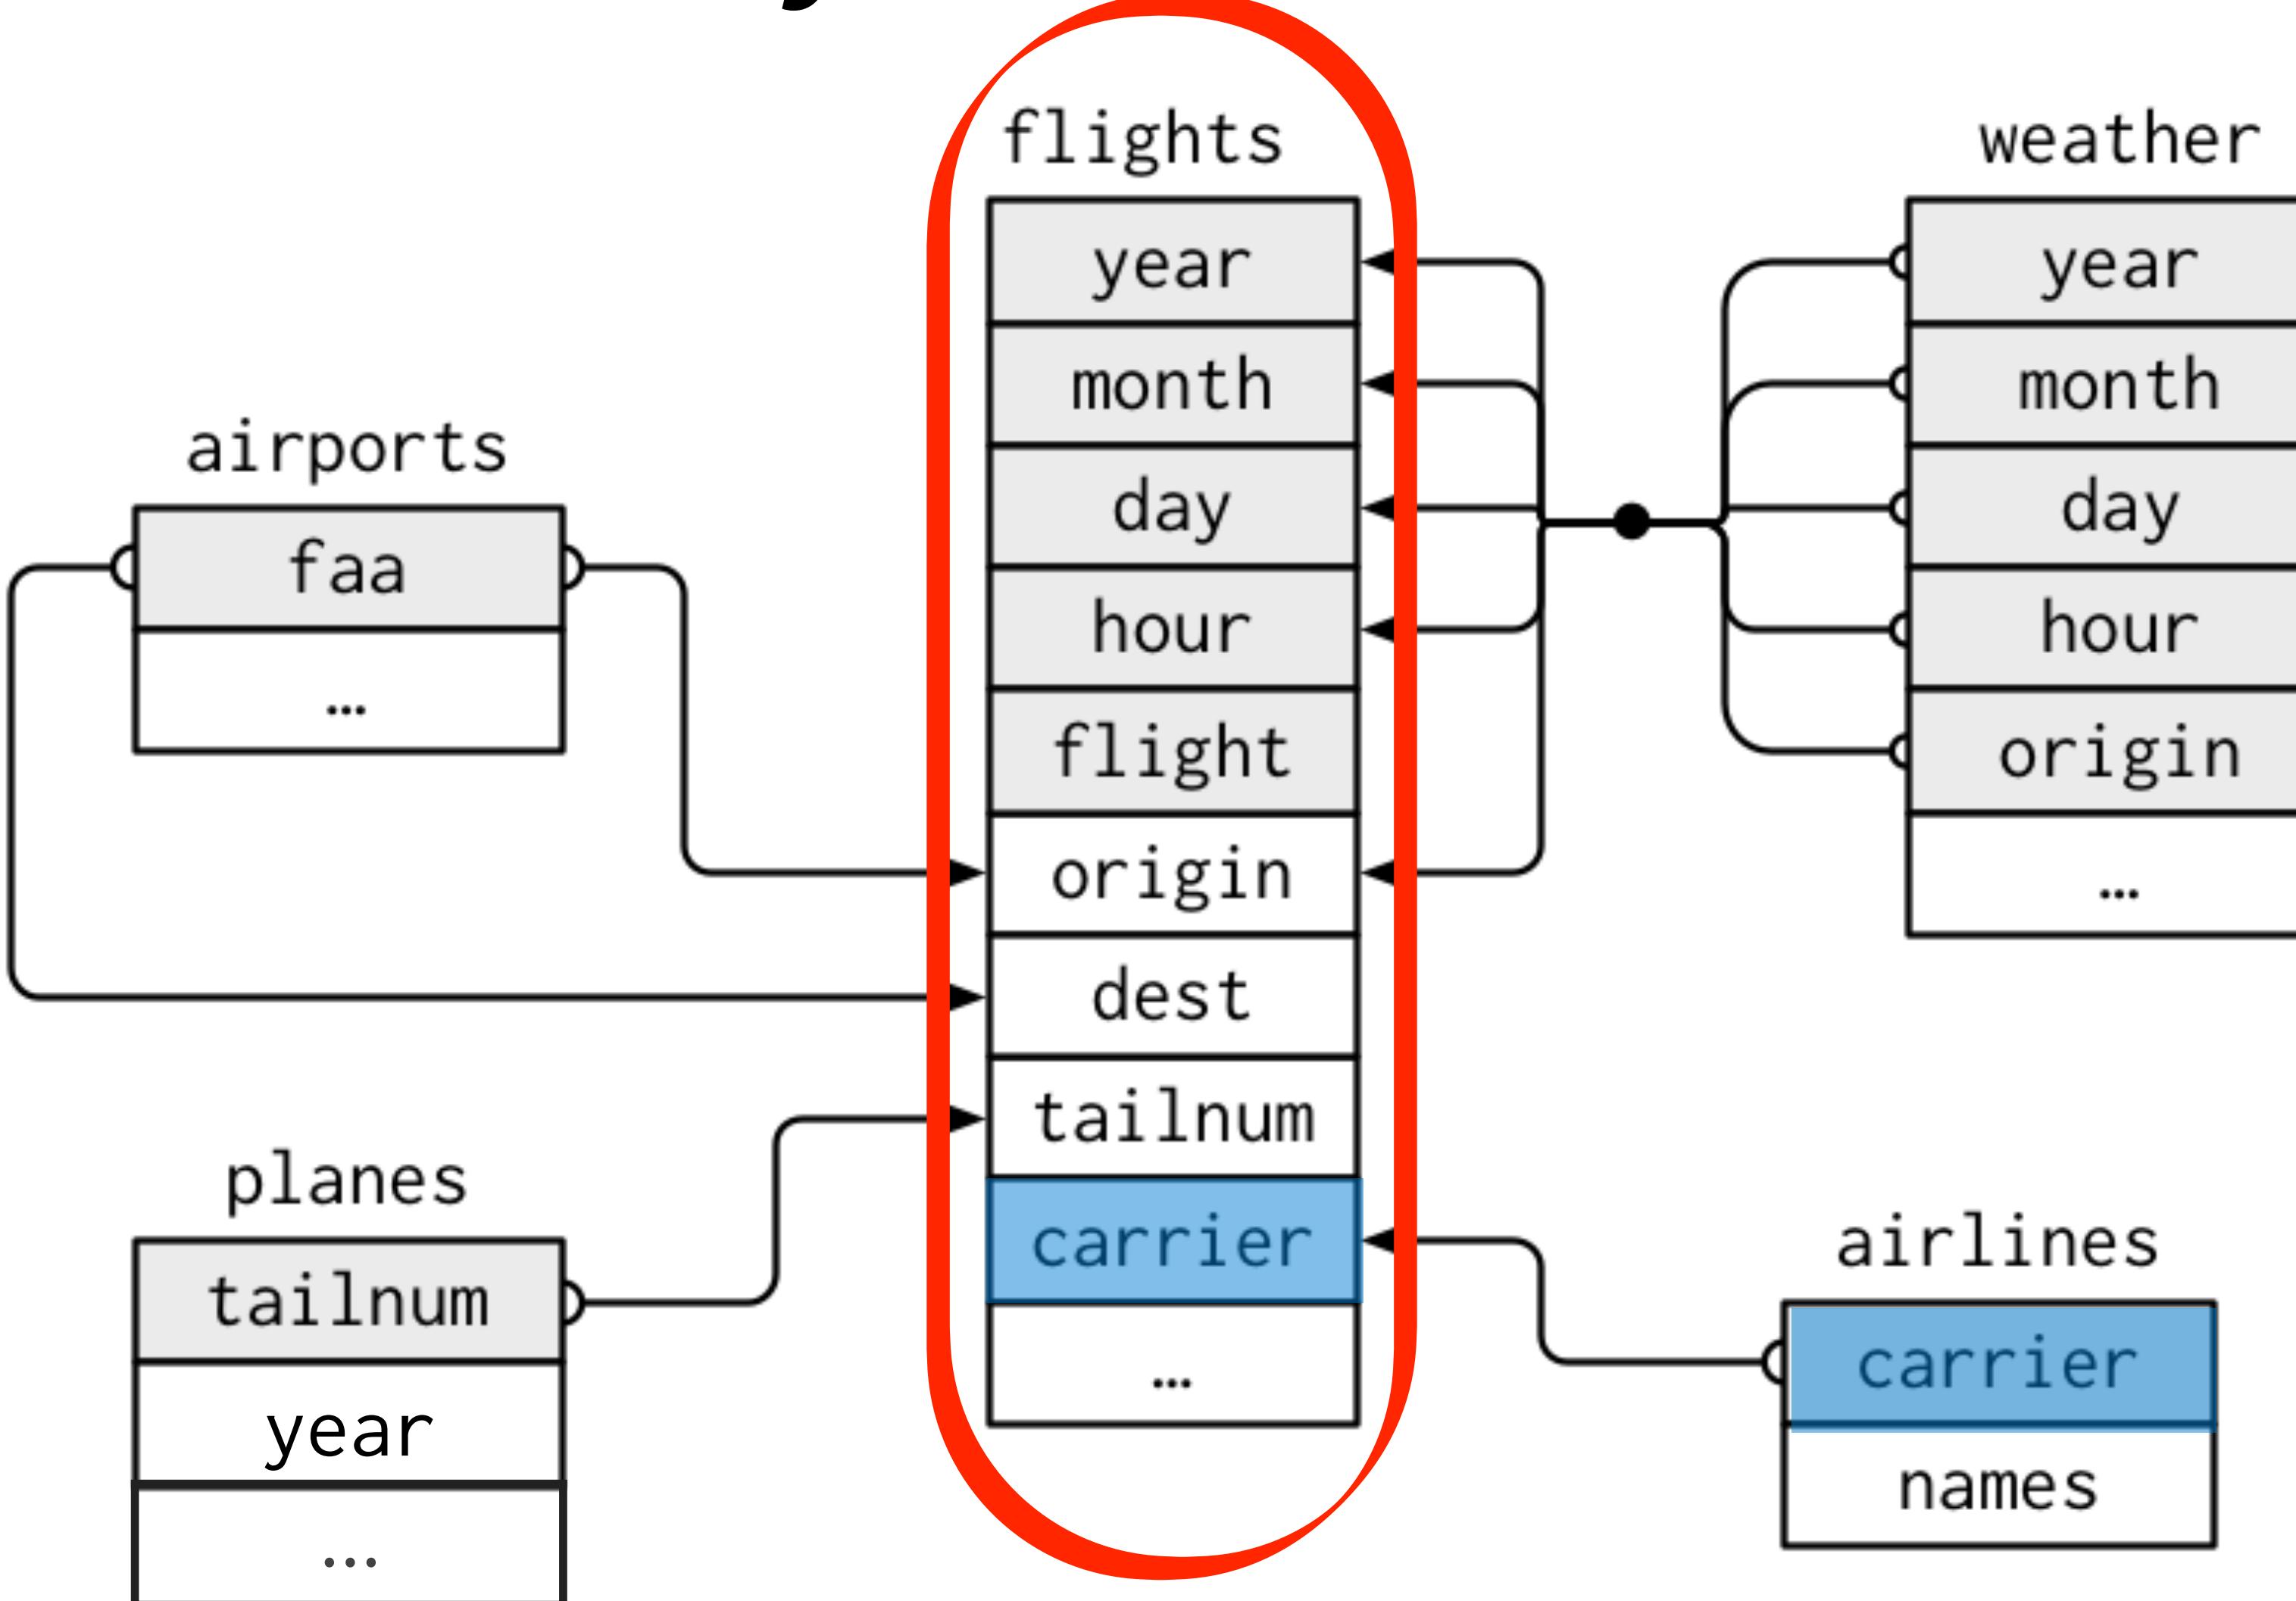
nycflights13



Data about every flight that departed La
Guardia, JFK, or Newark airports in 2013

```
# install.packages("nycflights13")
library(nycflights13)
```

nycflights13



Flights

[View\(flights\)](#)

year	month	day	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time	arr_delay	carrier	flight	tailnum	origin	dest	air_time	distance	hour	minute	time_hour
2013	1	1	517	515	2	830	819	11	UA	1545	N14228	EWR	IAH	227	1400	5	15	2013-01-01 05:00:00
2013	1	1	533	529	4	850	830	20	UA	1714	N24211	LGA	IAH	227	1416	5	29	2013-01-01 05:00:00
2013	1	1	542	540	2	923	850	33	AA	1141	N619AA	JFK	MIA	160	1089	5	40	2013-01-01 05:00:00
2013	1	1	544	545	-1	1004	1022	-18	B6	725	N804JB	JFK	BQN	183	1576	5	45	2013-01-01 05:00:00
2013	1	1	554	600	-6	812	837	-25	DL	461	N668DN	LGA	ATL	116	762	6	0	2013-01-01 06:00:00
2013	1	1	554	558	-4	740	728	12	UA	1696	N39463	EWR	ORD	150	719	5	58	2013-01-01 05:00:00
2013	1	1	555	600	-5	913	854	19	B6	507	N516JB	EWR	FLL	158	1065	6	0	2013-01-01 06:00:00
2013	1	1	557	600	-3	709	723	-14	EV	5708	N829AS	LGA	IAD	53	229	6	0	2013-01-01 06:00:00
2013	1	1	557	600	-3	838	846	-8	B6	79	N593JB	JFK	MCO	140	944	6	0	2013-01-01 06:00:00
2013	1	1	558	600	-2	753	745	8	AA	301	N3ALAA	LGA	ORD	138	733	6	0	2013-01-01 06:00:00
2013	1	1	558	600	-2	849	851	-2	B6	49	N793JB	JFK	PBI	149	1028	6	0	2013-01-01 06:00:00
2013	1	1	558	600	-2	853	856	-3	B6	71	N657JB	JFK	TPA	158	1005	6	0	2013-01-01 06:00:00
2013	1	1	558	600	-2	924	917	7	UA	194	N29129	JFK	LAX	345	2475	6	0	2013-01-01 06:00:00
2013	1	1	558	600	-2	923	937	-14	UA	1124	N53441	EWR	SFO	361	2565	6	0	2013-01-01 06:00:00
2013	1	1	559	600	-1	941	910	31	AA	707	N3DUAA	LGA	DFW	257	1389	6	0	2013-01-01 06:00:00
2013	1	1	559	559	0	702	706	-4	B6	1806	N708JB	JFK	BOS	44	187	5	59	2013-01-01 05:00:00
2013	1	1	559	600	-1	854	902	-8	UA	1187	N76515	EWR	LAS	337	2227	6	0	2013-01-01 06:00:00
2013	1	1	600	600	0	851	858	-7	B6	371	N595JB	LGA	FLL	152	1076	6	0	2013-01-01 06:00:00

Flights

What airlines have the longest delays?

carrier	avg_delay
9E	?
AA	?
AS	?
B6	?
DL	?
EV	?
F9	?
FL	?

name	avg_delay
AirTran Airways Corporation	?
Alaska Airlines Inc.	?
American Airlines Inc.	?
Delta Air Lines Inc.	?
Endeavor Air Inc.	?
Envoy Air	?
ExpressJet Airlines Inc.	?
Frontier Airlines Inc.	?

carrier	avg_delay
9E	?
AA	?
AS	?
B6	?
DL	?
EV	?
F9	?
FL	?

AIRLINES

FLIGHTS

name	avg_delay
AirTran Airways Corporation	?
Alaska Airlines Inc.	?
American Airlines Inc.	?
Delta Air Lines Inc.	?
Endeavor Air Inc.	?
Envoy Air	?
ExpressJet Airlines Inc.	?
Frontier Airlines Inc.	?

Airline names

[View\(flights\)](#)

[View\(airlines\)](#)

arr_delay	carrier
11	UA
20	UA
33	AA
-18	B6
-25	DL
12	UA

carrier	name
9E	Endeavor Air Inc.
AA	American Airlines Inc.
AS	Alaska Airlines Inc.
B6	JetBlue Airways
DL	Delta Air Lines Inc.
EV	ExpressJet Airlines Inc.

Airline names

[View\(flights\)](#)

arr_delay	carrier	name
11	UA	Endeavor Air Inc.
20	UA	American Airlines Inc.
33	AA	Alaska Airlines Inc.
-18	B6	JetBlue Airways
-25	DL	Delta Air Lines Inc.
12	UA	ExpressJet Airlines Inc.

Airline names

[View\(flights\)](#)

arr_delay	carrier	name
11	UA	Endeavor Air Inc.
20	UA	American Airlines Inc.
33	AA	Alaska Airlines Inc.
-18	B6	JetBlue Airways
-25	DL	Delta Air Lines Inc.
12	UA	ExpressJet Airlines Inc.

mutating joins

A faint watermark of the R logo is visible in the bottom right corner, consisting of a circular arrow and the letters "R".

common syntax

Each join function returns a data frame / tibble.

```
left_join(x, y, by = NULL, ...)
```

join function

data frames
to join

names of columns
to join on

Toy data

```
band <- tribble(  
  ~name,      ~band,  
  "Mick",    "Stones",  
  "John",    "Beatles",  
  "Paul",    "Beatles"  
)
```

band

name	band
Mick	Stones
John	Beatles
Paul	Beatles

```
instrument <- tribble(  
  ~name,    ~plays,  
  "John",   "guitar",  
  "Paul",   "bass",  
  "Keith",  "guitar"  
)
```

instrument

name	plays
John	guitar
Paul	bass
Keith	guitar

Toy data

band		instrument	
name	band	name	plays
Mick	Stones	John	guitar
John	Beatles	Paul	bass
Paul	Beatles	Keith	guitar

left

```
band %>% left_join(instrument, by = "name")
```

band

name	band
Mick	Stones
John	Beatles
Paul	Beatles

+

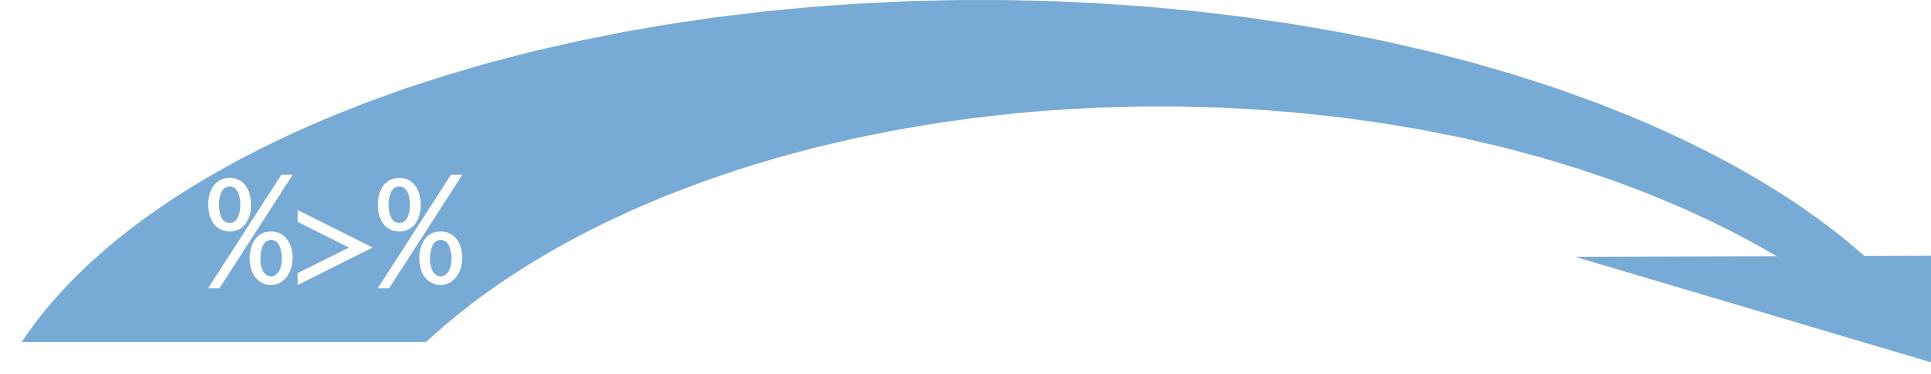
instrument

name	plays
John	guitar
Paul	bass

=

name	band	plays
Mick	Stones	<NA>
John	Beatles	guitar
Paul	Beatles	bass

The pipe operator %>%



%>%

instrument %>% left_join(_____, by = "name")

Passes result on left into first argument of function on right.
So, for example, these do the same thing. Try it.

```
left_join(instrument, by = "name")  
instrument %>% left_join(by = "name")
```

left

```
band %>% left_join(instrument, by = "name")
```

band

name	band
Mick	Stones
John	Beatles
Paul	Beatles

+

instrument

name	plays
John	guitar
Paul	bass

=

name	band	plays
Mick	Stones	<NA>
John	Beatles	guitar
Paul	Beatles	bass

right

```
band %>% right_join(instrument, by = "name")
```

band

name	band
Mick	Stones
John	Beatles
Paul	Beatles

+

instrument

name	plays
John	guitar
Paul	bass
Keith	guitar

=

name	band	plays
John	Beatles	guitar
Paul	Beatles	bass
Keith	<NA>	guitar

full

```
band %>% full_join(instrument, by = "name")
```

band

name	band
Mick	Stones
John	Beatles
Paul	Beatles

+

instrument

name	plays
John	guitar
Paul	bass
Keith	guitar

=

name	band	plays
Mick	Stones	<NA>
John	Beatles	guitar
Paul	Beatles	bass
Keith	<NA>	guitar

inner

```
band %>% inner_join(instrument, by = "name")
```

band

name	band
Mick	Stones
John	Beatles
Paul	Beatles

+

instrument

name	plays
John	guitar
Paul	bass
Keith	guitar

=

name	band	plays
John	Beatles	guitar
Paul	Beatles	bass

Airline names

[View\(flights\)](#)

[View\(airlines\)](#)

arr_delay	carrier
11	UA
20	UA
33	AA
-18	B6
-25	DL
12	UA

carrier	name
9E	Endeavor Air Inc.
AA	American Airlines Inc.
AS	Alaska Airlines Inc.
B6	JetBlue Airways
DL	Delta Air Lines Inc.
EV	ExpressJet Airlines Inc.

Your Turn 1

Which airlines had the largest arrival delays? Work in groups to complete the code below.

```
flights %>%  
filter(!is.na(arr_delay)) %>%  
_____ %>%  
group_by(_____) %>%  
_____ %>%  
arrange(____)
```

1. Join airlines to flights

2. Compute and order the average arrival delays by airline. Display full names, no codes.



```
flights %>%  
  filter(!is.na(arr_delay)) %>%  
  left_join(airlines, by = "carrier") %>%  
  group_by(name) %>%  
  summarise(delay = mean(arr_delay)) %>%  
  arrange(delay)  
  
## # A tibble: 16 × 2  
##   name          delay  
##   <chr>        <dbl>  
## 1 Alaska Airlines Inc. -9.9308886  
## 2 Hawaiian Airlines Inc. -6.9152047  
## 3 American Airlines Inc.  0.3642909  
## 4 Delta Air Lines Inc.   1.6443409  
## 5 Virgin America          1.7644644
```



Toy data

```
band <- tribble(  
  ~name,      ~band,  
  "Mick",    "Stones",  
  "John",    "Beatles",  
  "Paul",    "Beatles"  
)
```

band

name	band
Mick	Stones
John	Beatles
Paul	Beatles

```
instrument2 <- tribble(  
  ~artist,    ~plays,  
  "John",    "guitar",  
  "Paul",    "bass",  
  "Keith",   "guitar"  
)
```

instrument2

artist	plays
John	guitar
Paul	bass
Keith	guitar

What if the names do not match?

Use a named vector to match on variables with different names.

```
band %>% left_join(instrument2, by = c("name" = "artist"))
```

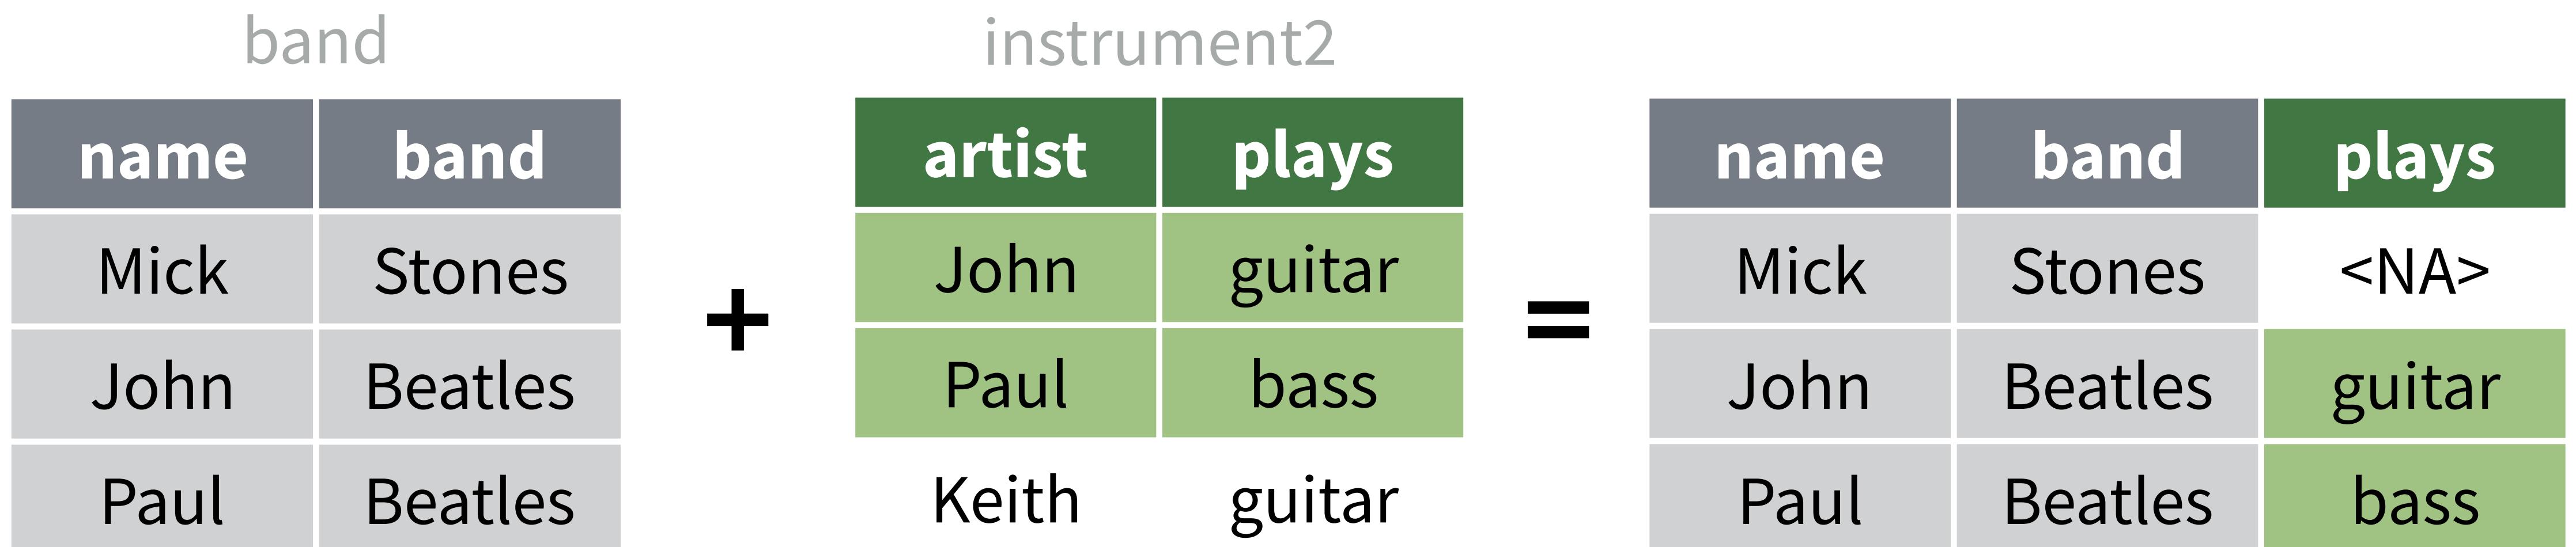
A named vector

The name of the
element = the column
name in the first data
set

The value of the
element = the column
name in the second
data set

common syntax - matching names

```
band %>% left_join(instrument2, by = c("name" = "artist"))
```



Airport names

`View(flights["dest"])`

	dest
1	IAH
2	IAH
3	MIA
4	BQN
5	ATL
6	ORD
7	FIA

`View(airports)`

faa	name
04G	Lansdowne Airport
06A	Moton Field Municipal Airport
06C	Schaumburg Regional
06N	Randall Airport
09J	Jekyll Island Airport
0A9	Elizabethton Municipal Airport
0CC	Willow Creek Airport

Airport names

```
airports %>% select(1:3)
```

	faa	name
	<chr>	<chr>
04G	Lansdowne Airport	
06A	Moton Field Municipal Airport	
06C	Schaumburg Regional	
06N	Randall Airport	
09J	Jekyll Island Airport	
0A9	Elizabethton Municipal Airport	
0G6	Williams County Airport	
0G7	Finger Lakes Regional Airport	

```
flights %>% select(14:15)
```

	dest	air_time
	<chr>	<dbl>
	IAH	227
	IAH	227
	MIA	160
	BQN	183
	ATL	116
	ORD	150
	FLL	158
	IAD	53

common syntax - matching names

```
airports %>% left_join(flights, by = c("faa" = "dest"))
```

	faa <chr>	name <chr>		dest <chr>	air_time <dbl>
04G	Lansdowne Airport			IAH	227
06A	Moton Field Municipal Airport			IAH	227
06C	Schaumburg Regional			MIA	160
06N	Randall Airport			BQN	183
09J	Jekyll Island Airport			ATL	116
0A9	Elizabethton Municipal Airport			ORD	150
0G6	Williams County Airport			FLL	158
0G7	Finger Lakes Regional Airport			IAD	53

Your Turn 2

Use flights and airports to compute the distance and average arr_delay by destination airport (names only, not codes). Order by average delay, worst to best.

Hint: use first() to get distance.



```
flights %>%  
  filter(!is.na(arr_delay)) %>%  
  left_join(airports, by = c("dest" = "faa")) %>%  
  group_by(name) %>%  
  summarise(distance = first(distance),  
            delay = mean(arr_delay)) %>%  
  arrange(desc(delay))  
## # A tibble: 101 × 3  
## # ... with 3 variables:  
## #   name <chr>    distance <dbl>    delay <dbl>  
## #   1 Columbia Metropolitan      602 41.76415  
## #   2 Tulsa Intl                 1215 33.65986  
## #   3 Will Rogers World          1325 30.61905
```

filtering joins

A faint watermark of the R logo is visible in the bottom right corner of the slide.

semi

```
band %>% semi_join(instrument, by = "name")
```

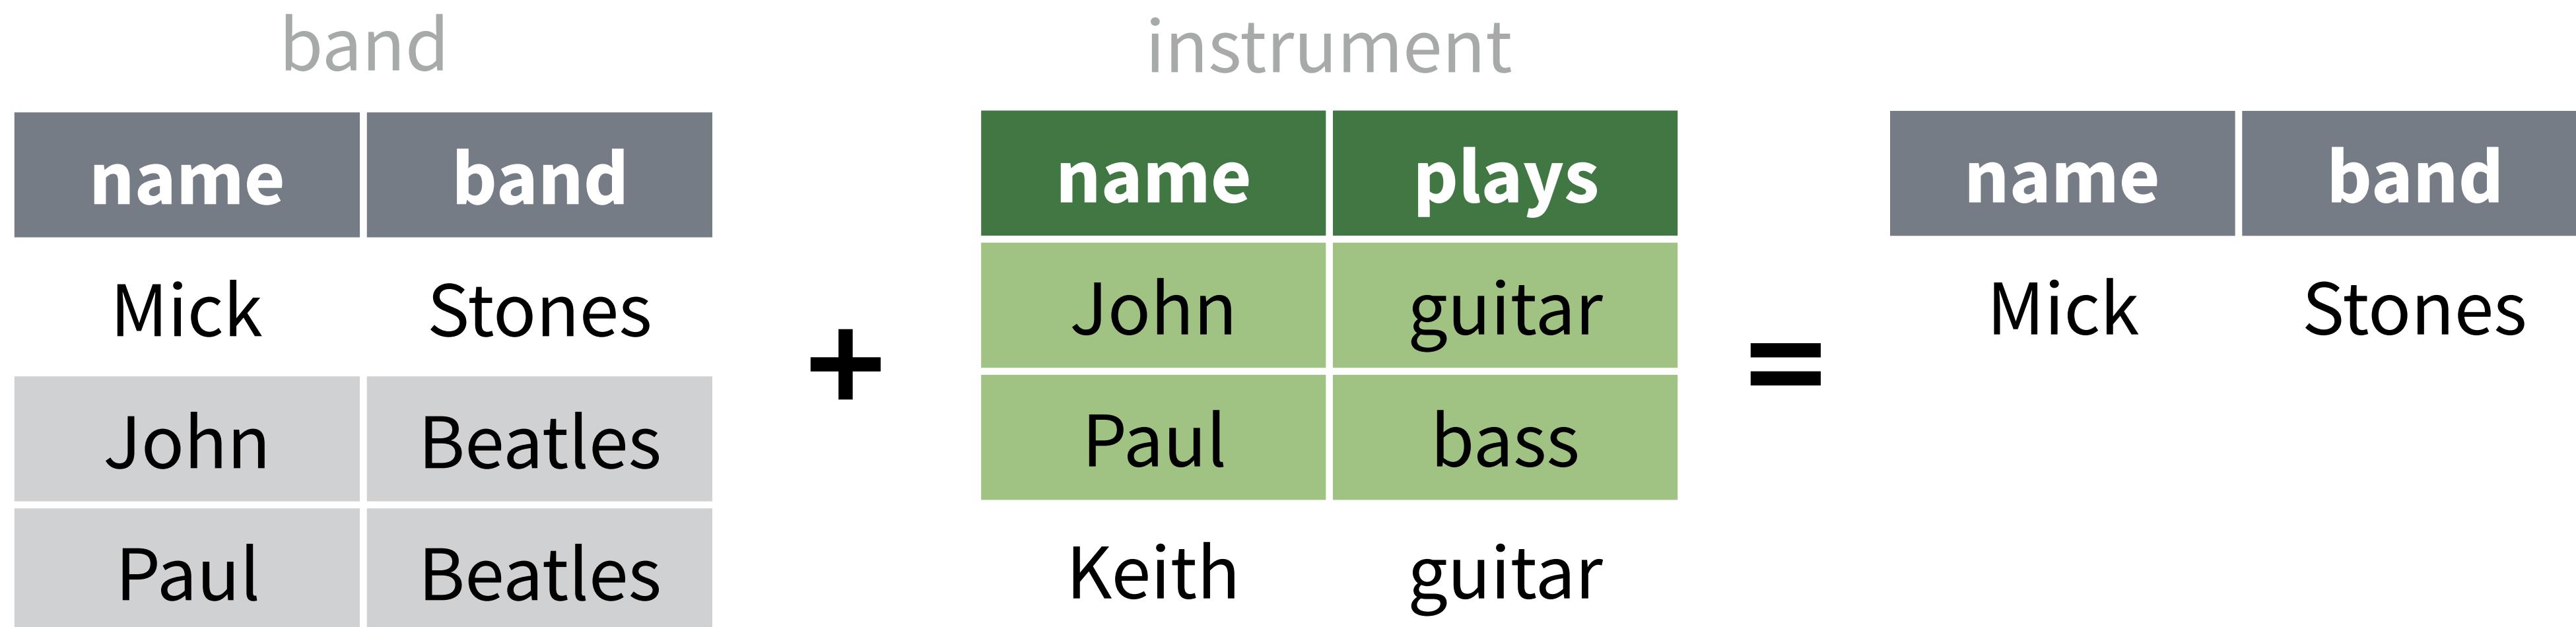
band		instrument			
name	band	name	plays	name	band
Mick	Stones	John	guitar	John	Beatles
John	Beatles	Paul	bass	Paul	Beatles
Paul	Beatles	Keith	guitar		

+

=

anti

```
band %>% anti_join(instrument, by = "name")
```



Airport names

```
airports %>% select(1:3)
```

	faa	name
	<chr>	<chr>
04G	Lansdowne Airport	
06A	Moton Field Municipal Airport	
06C	Schaumburg Regional	
06N	Randall Airport	
09J	Jekyll Island Airport	
0A9	Elizabethton Municipal Airport	
0G6	Williams County Airport	
0G7	Finger Lakes Regional Airport	

```
flights %>% select(14:15)
```

	dest	air_time
	<chr>	<dbl>
	IAH	227
	IAH	227
	MIA	160
	BQN	183
	ATL	116
	ORD	150
	FLL	158
	IAD	53

Your Turn 3

How many airports in **airports** are serviced by flights originating in New York (i.e. flights in our dataset?)

Notice that the column to join on is named **faa** in the **airports** dataset and **dest** in the **flights** dataset.



```
airports %>%  
  semi_join(flights, by = c("faa" = "dest")) %>%  
  distinct(faa)
```

faa
<chr>
IAH
MIA
ATL
ORD
FLL
IAD
MCO
PBI
TPA
LAX

1-10 of 101 rows

Previous 1 2 3 4 5 6 ... 11 Next



distinct()

Removes rows with duplicate values (in a column).

```
distinct(instrument, plays)
```

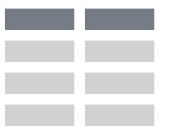
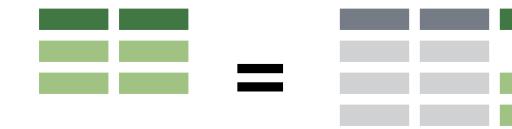
instrument

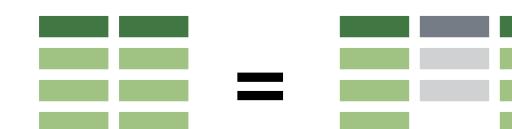
name	plays
John	guitar
Paul	bass
Keith	guitar

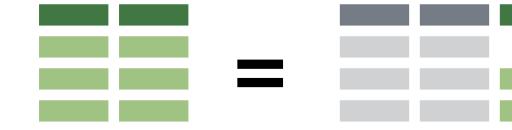
→

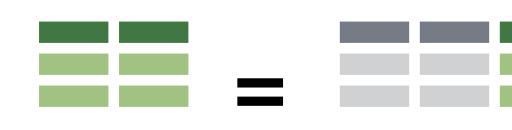
plays
guitar
bass

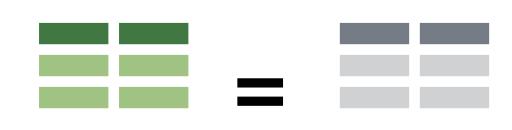
Recap: Two table verbs

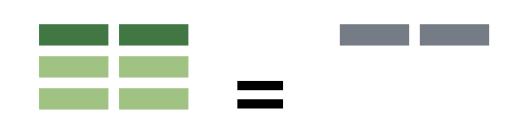
 +  =  **left_join()** retains all cases in **left** data set

 +  =  **right_join()** retains all cases in **right** data set

 +  =  **full_join()** retains all cases in **either** data set

 +  =  **inner_join()** retains only cases in **both** data sets

 +  =  **semi_join()** extracts cases that **have a match**

 +  =  **anti_join()** extracts cases that **do not have a match**



Two table verbs

Data Transformation cheatsheet

ON BACK

Vector Functions

Summary Functions

Combine Tables

dplyr

This section contains a cheatsheet for dplyr's two-table verbs. It includes sections for Vector Functions, Summary Functions, and Combine Tables. The Vector Functions and Summary Functions sections provide details on how to use them with mutate() and summarise(). The Combine Tables section provides examples of how to use bind_cols(), bind_rows(), left_join(), right_join(), inner_join(), and full_join().

Combine Tables

COMBINE VARIABLES

x	y	=
A B C a t 1 b u 2 c v 3	A B D a t 3 b u 2 d w 1	A B C A B D a t 1 a t 3 b u 2 b u 2 c v 3 d w 1

Use `bind_cols()` to paste tables beside each other as they are.

`bind_cols(...)` Returns tables placed side by side as a single table.
BE SURE THAT ROWS ALIGN.

Use a "Mutating Join" to join one table to columns from another, matching values with the rows that they correspond to. Each join retains a different combination of values from the tables.

`left_join(x, y, by = NULL, copy = FALSE, suffix = c(".x", ".y"), ...)`
Join matching values from y to x.

`right_join(x, y, by = NULL, copy = FALSE, suffix = c(".x", ".y"), ...)`
Join matching values from x to y.

`inner_join(x, y, by = NULL, copy = FALSE, suffix = c(".x", ".y"), ...)`
Join data. Retain only rows with matches.

`full_join(x, y, by = NULL, copy = FALSE, suffix = c(".x", ".y"), ...)`
Join data. Retain all values, all rows.

Use `by = c("col1", "col2", ...)` to specify one or more common columns to match on.
`left_join(x, y, by = "A")`

Use a named vector, `by = c("col1" = "col2")`, to match on columns that have different names in each table.
`left_join(x, y, by = c("C" = "D"))`

Use `suffix` to specify the suffix to give to unmatched columns that have the same name in both tables.
`left_join(x, y, by = c("C" = "D"), suffix = c("1", "2"))`

COMBINE CASES

x	y	=
A B C a t 1 b u 2 c v 3	A B C C v 3 d w 4	A B C a t 1 b u 2 c v 3 C v 3 d w 4

Use `bind_rows()` to paste tables below each other as they are.

`bind_rows(..., .id = NULL)`
Returns tables one on top of the other as a single table. Set `.id` to a column name to add a column of the original table names (as pictured)

`intersect(x, y, ...)`
Rows that appear in both x and y.

`setdiff(x, y, ...)`
Rows that appear in x but not y.

`union(x, y, ...)`
Rows that appear in x or y.
(Duplicates removed). `union_all()` retains duplicates.

Use `setequal()` to test whether two data sets contain the exact same rows (in any order).

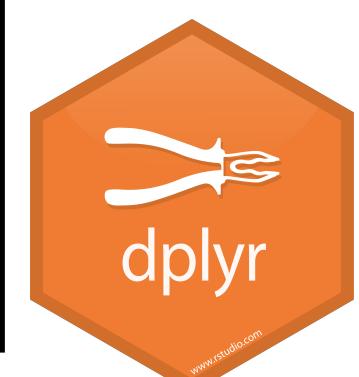
EXTRACT ROWS

x	y	=
A B C a t 1 b u 2 c v 3	A B D a t 3 b u 2 d w 1	A B C a t 3 b u 2 d w 1

Use a "Filtering Join" to filter one table against the rows of another.

`semi_join(x, y, by = NULL, ...)`
Return rows of x that have a match in y.
USEFUL TO SEE WHAT WILL BE JOINED.

`anti_join(x, y, by = NULL, ...)`
Return rows of x that do not have a match in y. USEFUL TO SEE WHAT WILL NOT BE JOINED.



Join Data with

